

Implementation of iLQR to Solve Gymnasium Tasks

Ganga Nair B
M.Tech RAS, 1st Year

E2 335 – Topics in Artificial Intelligence, 2025

- 1 Motivation and Introduction
- 2 Aim and Objective
- 3 Scope
- 4 Theory Behind iLQR
- 5 Literature Review
- 6 Implementation on Mountain Car
- 7 Model-Free Extension using Neural Networks
- 8 Application to Swing-Up Cart Pole
- 9 Results and Conclusion
- 10 Future Work

- Reinforcement Learning (RL) aims to find optimal policies maximizing cumulative reward.
- Optimal control seeks actions that minimize cost while achieving goals.
- RL has deep roots in control theory.
- One such method is the Iterative Linear Quadratic Regulator (iLQR).

- iLQR provides near-optimal control policies analytically with fewer iterations.
- Offers theoretical guarantees, interpretability, and low computational cost.
- Assumes accurate dynamics of the system are known.

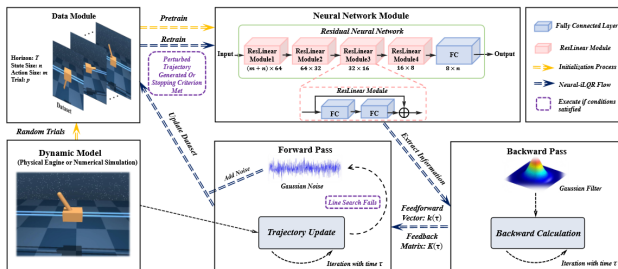
To solve benchmark control problems in the Gymnasium environment using model-based RL methods, and extend the framework to unknown system dynamics.

- Apply iLQR to classical Gym tasks:
 - Mountain Car
 - Swing-Up Cart Pole
- Learn dynamics using a neural network from scratch.
- Use learned dynamics in iLQR to control the system.

Literature Review: Neural-iLQR & Related Works

Neural-iLQR: Trajectory Optimization

- **Objective:** Overcome model inaccuracies in conventional iLQR for nonlinear systems
- **Solution:**
 - Integrates 2 layered NN to learn local dynamics online
 - Alternates between policy optimization & network training



Physics-Informed Vehicle Tracking

- **Objective:** Improve trajectory tracking accuracy with physical consistency
- **Solution:**
 - Combines PINN with Kinematic Bicycle Model
 - ILQR controller used for control

What is Linear Quadratic Regulator (LQR)?

- Solves constrained optimization using dynamic programming.
- Backward pass computes gain matrices K and k .
- Forward pass generates control trajectory using the optimized gains.

Linear Quadratic Regulator (LQR)

Assumptions:

- Linear time-varying system dynamics:

$$\mathbf{x}_{t+1} = \mathbf{F}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \mathbf{f}_t$$

- Quadratic cost function:

$$c(\mathbf{x}_t, \mathbf{u}_t) = \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{C}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{c}_t$$

LQR Problem Statement

Objective:

$$\min_{\mathbf{u}_0, \dots, \mathbf{u}_{T-1}} \sum_{t=0}^{T-1} c(\mathbf{x}_t, \mathbf{u}_t)$$

Subject to:

$$\mathbf{x}_{t+1} = \mathbf{F}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \mathbf{f}_t$$

LQR with Dynamic Programming

- **Cost optimization:** Minimize $\sum_{t=0}^{T-1} c(x_t, u_t)$ iteratively (Bellman-style).
- Expressed via **Value** (V) and **Action-Value** (Q) functions:

Q-Function Definition

$$Q_t = \text{Current cost} + V(\text{Next step})$$

$$Q(x_{T-1}, u_{T-1}) = \text{const} + \frac{1}{2} \begin{bmatrix} x_{T-1} \\ u_{T-1} \end{bmatrix}^T C_{T-1} \begin{bmatrix} x_{T-1} \\ u_{T-1} \end{bmatrix} + \begin{bmatrix} x_{T-1} \\ u_{T-1} \end{bmatrix}^T c_{T-1} + V(x_T)$$

- **Key idea:** Solve backward from $T - 1$ to 0 using dynamic programming.

- **Iterate from last step to first step**

- 1 Substitute Q_{T-1} with V_T expressed in terms of x_T (where $x_T = Ax_{T-1} + Bu_{T-1}$).
- 2 Optimize Q by solving $\nabla Q_{T-1} = 0$.
- 3 Compute feedback gains \mathbf{K}_{T-1} and \mathbf{k}_{T-1} to get control:

$$u_{T-1} = \mathbf{K}_{T-1}x_{T-1} + \mathbf{k}_{T-1}$$

- 4 Substitute u_{T-1} back into Q_{T-1} to obtain V_{T-1} (used for step $T - 2$).

iLQR Backward Pass: Algorithm Overview

Initialization (Terminal Step)

$$V_x = Q_{\text{terminal}}(x_T - x_{\text{target}}), \quad V_{xx} = Q_{\text{terminal}}$$

Action-Value Function $Q(x_t, u_t)$

$$Q(x_t, u_t) = \underbrace{\frac{1}{2} \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T \begin{bmatrix} Q_{xx} & Q_{xu} \\ Q_{ux} & Q_{uu} \end{bmatrix} \begin{bmatrix} x_t \\ u_t \end{bmatrix}}_{\text{Quadratic term}} + \underbrace{\begin{bmatrix} Q_x & Q_u \end{bmatrix} \begin{bmatrix} x_t \\ u_t \end{bmatrix}}_{\text{Linear term}} + \text{const}$$

Value Function $V(x_t)$ (After Optimization)

$$V(x_t) = \underbrace{\frac{1}{2} x_t^T V_{xx} x_t}_{\text{Quadratic term}} + \underbrace{V_x^T x_t}_{\text{Linear term}} + \text{const}$$

LQR Backward Pass: Optimization

Key Quantities

$$Q_x = A_t^\top V_x + C_x, \quad Q_u = B_t^\top V_x + Ru_t$$

$$Q_{xx} = A_t^\top V_{xx} A_t + C_{xx}, \quad Q_{uu} = B_t^\top V_{xx} B_t + R$$

$$Q_{ux} = B_t^\top V_{xx} A_t + C_{ux}$$

Control Update

Solve $\nabla_u Q = 0$ to get feedback law:

$$k_t = -Q_{uu}^{-1} Q_u, \quad K_t = -Q_{uu}^{-1} Q_{ux}$$

Value Function Update

$$V_x \leftarrow Q_x + K_t^\top Q_{uu} k_t + K_t^\top Q_u$$

$$V_{xx} \leftarrow Q_{xx} + K_t^\top Q_{uu} K_t + K_t^\top Q_{ux} + Q_{ux}^\top K_t$$

- **Output:** Gain matrices K_t and offsets k_t for all t .

- Instead of pretending your nonlinear system is globally linear, you locally linearize it around a “guess” trajectory, solve an LQR problem there, then update your guess and repeat.

iLQR Mathematics: Nonlinear System Approximation

Nonlinear System Assumptions

- Dynamics: $x_{t+1} = f(x_t, u_t)$
- Cost: $c_t = \ell(x_t, u_t)$

Local Approximations (1st/2nd Order)

- **Dynamics (1st Order):**

$$\delta x_{t+1} = \underbrace{\nabla_{x_t} f(\hat{x}_t, \hat{u}_t)}_{A_t} \delta x_t + \underbrace{\nabla_{u_t} f(\hat{x}_t, \hat{u}_t)}_{B_t} \delta u_t$$

where $\delta x_t = x_t - \hat{x}_t$, $\delta u_t = u_t - \hat{u}_t$

- **Cost (2nd Order):**

$$\delta c_t = \underbrace{\nabla \ell(\hat{x}_t, \hat{u}_t)}_{c_t} \begin{bmatrix} \delta x_t \\ \delta u_t \end{bmatrix} + \frac{1}{2} \begin{bmatrix} \delta x_t \\ \delta u_t \end{bmatrix}^T \underbrace{\nabla^2 \ell(\hat{x}_t, \hat{u}_t)}_{C_t} \begin{bmatrix} \delta x_t \\ \delta u_t \end{bmatrix}$$

Iterative Linear Quadratic Regulator (iLQR)

Motivation:

- LQR assumes linear dynamics; not suitable for nonlinear systems.
- iLQR extends LQR to handle nonlinear dynamics through iteration.

Approach:

- 1 Initialize with a nominal trajectory.
- 2 Linearize dynamics and quadratize cost around the nominal trajectory.
- 3 Apply LQR to compute control updates.
- 4 Update nominal trajectory and repeat until convergence.

Iterative LQR (iLQR) Algorithm

WHILE not converged

1. Linearization (Current Trajectory \hat{x}_t, \hat{u}_t):

- Dynamics Jacobian:

$$F_t = \nabla_{x_t, u_t} f(\hat{x}_t, \hat{u}_t) = \begin{bmatrix} A_t & B_t \end{bmatrix}$$

- Cost gradients/Hessians:

$$c_t = \nabla_{x_t, u_t} c(\hat{x}_t, \hat{u}_t), \quad C_t = \nabla_{x_t, u_t}^2 c(\hat{x}_t, \hat{u}_t) = \begin{bmatrix} Q_{xx} & Q_{xu} \\ Q_{ux} & Q_{uu} \end{bmatrix}$$

2. Backward Pass (LQR on Deviations):

- Solve for $\delta u_t = K_t \delta x_t + k_t$ where:

$$\delta x_t = x_t - \hat{x}_t, \quad \delta u_t = u_t - \hat{u}_t$$

3. Forward Pass (Nonlinear Rollout):

- Apply control with feedback:

$$u_t = K_t(x_t - \hat{x}_t) + k_t + \hat{u}_t$$

- Update nominal trajectory (\hat{x}_t, \hat{u}_t) from rollout.

What If Dynamics Are Unknown?

- Estimate dynamics using a neural network.
- Train NN with randomly collected data.
- Improve the model iteratively as more data is collected.


Proposed Improvement:

- **Data Collection:** Generate dataset from iLQR-based rollouts, storing transitions (s_t, a_t, r_t, s_{t+1}) .
- **Neural Network for System Dynamics:** Train NN to predict next state: $s_{t+1} = f_{\theta}(s_t, a_t)$.
- **Control Optimization:** Use learned NN in iLQR instead of true model.

Neural Networks

- A **Neural Network (NN)** acts as a function approximator here.
- It maps inputs \mathbf{x} to outputs \mathbf{y} via layers of *neurons*.

Basic Structure:



- **Input Layer:** Receives state/control vectors.
- **Hidden Layers:** Perform successive linear transforms + nonlinear

Training Neural Networks

- ➊ **Data Collection:** Gather state–action–next-state tuples (x_t, u_t, x_{t+1}) .
- ➋ **Loss Function:** Measure prediction error, e.g.

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \|x_{t+1}^{(i)} - \hat{f}_{\theta}(x_t^{(i)}, u_t^{(i)})\|^2.$$

- ➌ **Optimization:** Use backpropagation + Optimization (like SGD) to adjust weights θ .
- ➍ **Regularization:** Techniques like dropout, weight decay to avoid overfitting.

Why Use Neural Networks to Estimate Dynamics for iLQR

- **Expressiveness:** Can approximate complex, nonlinear system dynamics better than fixed polynomial functions.
- **Data Efficiency:** Learns from observed transitions, capturing unmodeled effects (friction, aerodynamic drag).
- **Analytic Derivatives:** Modern frameworks provide $\partial \hat{f} / \partial x$ and $\partial \hat{f} / \partial u$ via autodiff for local linearization.
- **Modular Integration:** Use the learned model in iLQR algorithm:
 - *Forward pass:* simulate with \hat{f}_θ .
 - *Backward pass:* compute Jacobians A_t, B_t for LQR update.

Implementation: Solving Gym Classic Control Tasks

- Evaluate our approach with implementation on Mountain Car.

Mountain Car Continuous Environment

Problem Description

- **Goal:** Drive underpowered car up steep mountain
- **Challenge:** Must build momentum to climb
- **State:** Position $x \in [-1.2, 0.6]$, Velocity $\dot{x} \in [-0.07, 0.07]$
- **Actions:** $a \in [-1, 1]$





Illustration of MountainCar environment

MountainCar Dynamics & Gym Simplifications



Real-World Physics (x-direction)

$$\sum F_x = \underbrace{F_{\text{engine}}}_{\text{control}} - \underbrace{mg \sin \theta}_{\text{gravity}} - \underbrace{\mu mg \cos \theta}_{\text{friction}}$$

Dynamics

- Physics model:

$$\dot{x}_{t+1} = \dot{x}_t + 0.001a_t - 0.0025 \cos(3x_t)$$

$$x_{t+1} = x_t + \dot{x}_{t+1}$$

- **Reward:** -1 per timestep until goal ($x \geq 0.5$)
 - **Termination:** Reach flag or 200 steps
-
- Nonlinear dynamics make LQR inapplicable.
 - Applied iLQR to control task.

Gym Implementation

Gym Environment Assumptions

- **Friction ignored:** $\mu = 0$
- **Slope relation:** $\theta = 3x_t$ $(y = \sin(3x))$
- **Simplified gravity term:** $mg \sin \theta \rightarrow 0.0025 \cos(3x_t)$
- **Engine acceleration:** $F_{\text{engine}}/m \rightarrow \text{action} \times 0.001$

$$v_{t+1} = v_t + \underbrace{\text{action} \times 0.001}_{\text{engine}} - \underbrace{0.0025 \cos(3x_t)}_{\text{gravity}}$$
$$x_{t+1} = x_t + v_{t+1}$$

Gravity term $\cos(3x_t)$ creates a complex energy landscape:

Harder to climb as x increases

Attempt 1: LQR Implementation on MountainCar

Approach

- Implements classic Linear Quadratic Regulator (LQR) control.
- Linearizes the nonlinear dynamics at each timestep:

$$x_{t+1} \approx Ax_t + Bu_t$$

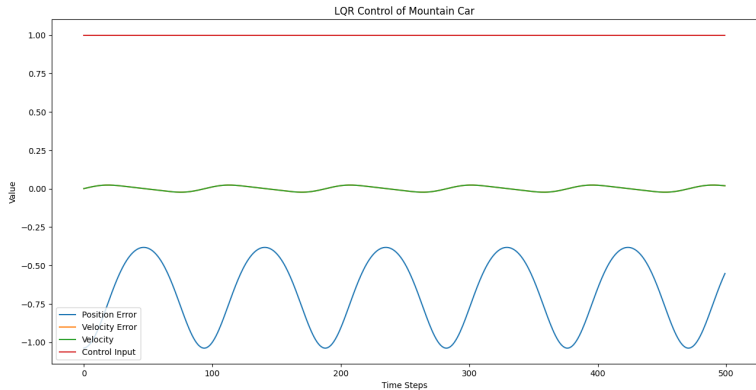
- Solves the discrete-time Riccati equation for feedback gain K .
- Control input:

$$u_t = -Kx_t$$

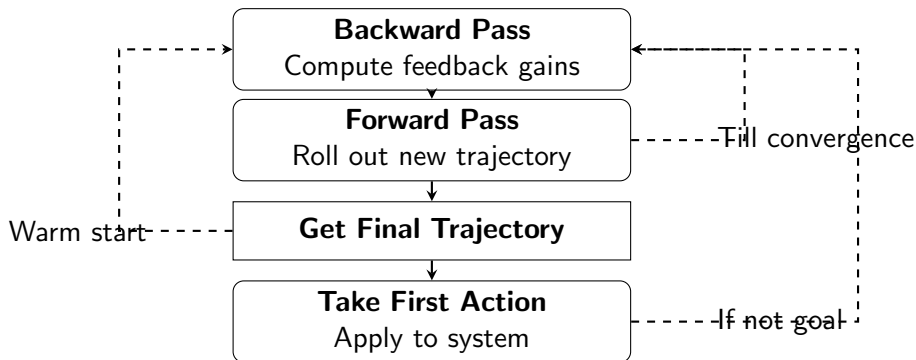
Limitations

- Assumes linear dynamics — MountainCar is inherently nonlinear.
- No consideration of $\cos(3x)$ term during linearization.
- Controller performs poorly near steep slopes.

Result



Iterative Control Loop Overview




Mountain Car implementation details

Hyperparameters:

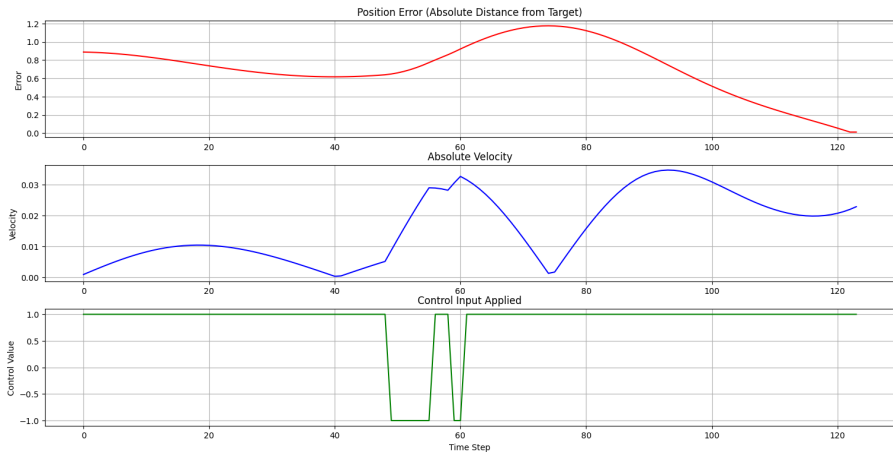
- Horizon: 40
- Constant number of iterations of LQR: 5
- $Q_{terminal} = \begin{bmatrix} 10 & 0 \\ 0 & 0 \end{bmatrix}$
- $R = [0.1]$
The known dynamics is provided.
- Backward pass uses linearised dynamics A, B
- Forward pass uses actual non linear dynamics

Results: iLQR with Known Model



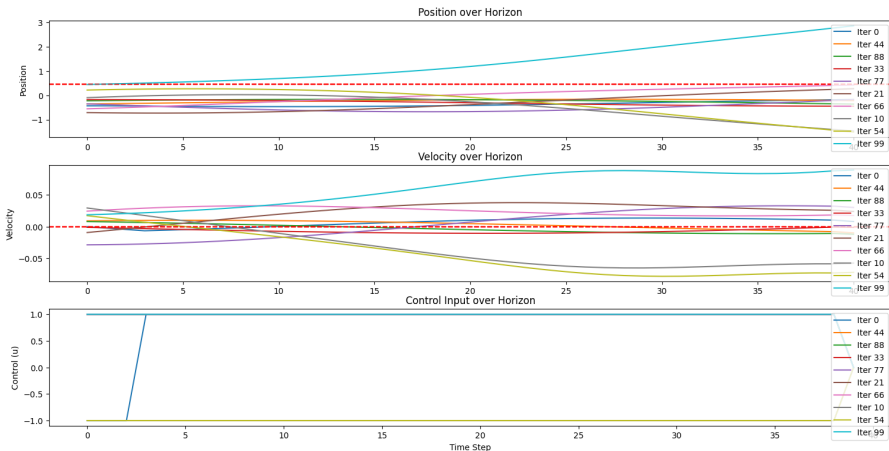
Rewards obtained and steps needed to reach goal over 50 iterations.

Results: iLQR with Known Model



Position error, Velocity and control action in one run

Results: iLQR with Known Model



10 different LQR optimised trajectories within the optimisation loop

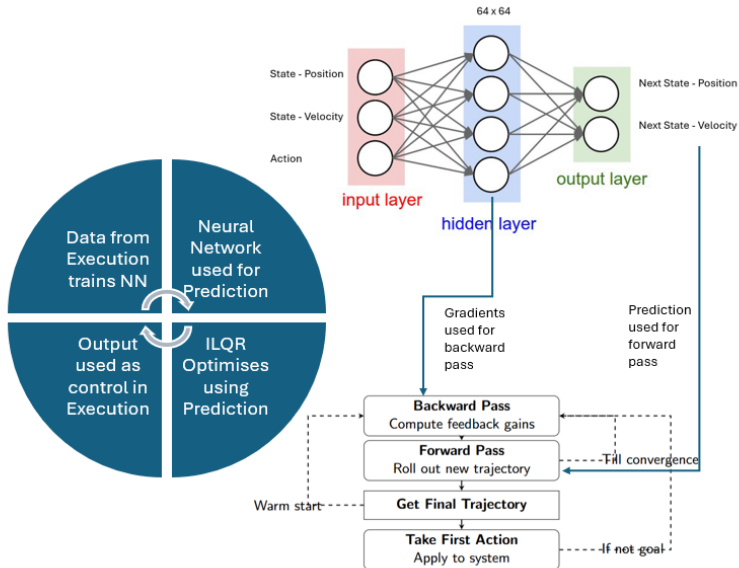
What if Dynamics are not known?

Using a Neural Network to predict dynamics

What if Dynamics are not known?

- Use a feedforward neural network to approximate unknown system dynamics.
- Architecture:
 - **Input:** State (x) and action (u)
 - **Network:** Two hidden layers with 64 ReLU units each
Input \rightarrow [64 neurons] \rightarrow [64 neurons] \rightarrow Output
 - **Output:** Next state prediction (x_{t+1})
- Model: MLPRegressor from sklearn
- Trained online using data collected from environment rollouts
- Mini-batch gradient descent using the Adam optimizer
- Initially random exploration phase helps initialise NN.

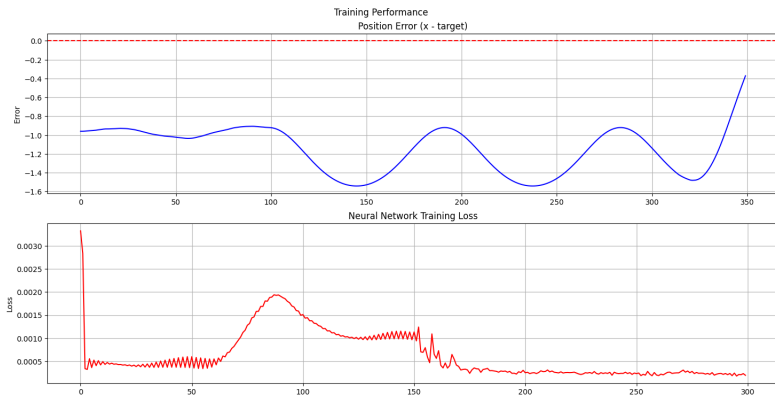
iLQR with NN implementation



Implementation Details for Neural Network

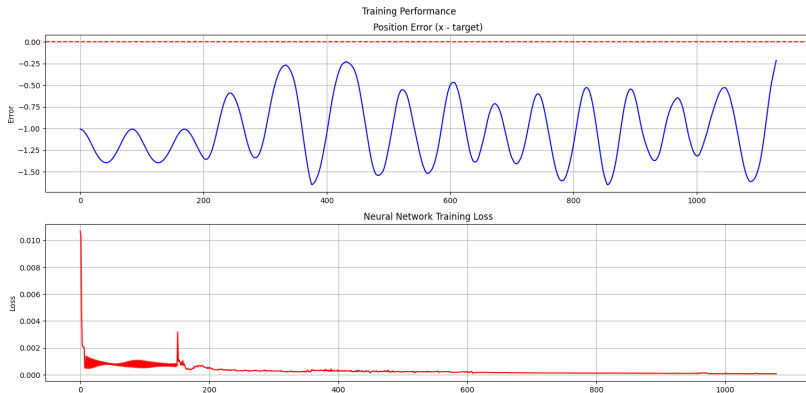
- ① **Random Exploration Phase** to initialize the network:
 - Approximately 200 steps of random exploration.
 - To encourage exploration across the state space: Apply action = +1 if position > 0, else -1.
- ② **Planning Horizon**: Set to 50 time steps.
- ③ **Training Schedule**: The neural network is updated every 50 steps.
- ④ **Gradient Estimation for Linearization**: Finite differences used to estimate Jacobians:
 - $A[:, i] = \frac{f(x+\varepsilon_i, u) - f(x, u)}{\varepsilon}$
 - Approximates $\partial f / \partial x$ and $\partial f / \partial u$ for use in iLQR backward pass.

Results: iLQR with NN



Results from partially integrated iLQR with gradients still obtained from known model

Results: iLQR with NN




Results from model in which known model completely removed

Swing-Up CartPole Environment

Problem Description

- **Goal:** Swing and balance the pole in upright position on a moving cart.
- **Challenge:** Starts in downward position; must swing up and stabilize.
- **State:** Position x , Velocity \dot{x} , Angle θ , Angular Velocity $\dot{\theta}$
- **Actions:** Continuous force a in $[-10, 10]$



Gym environment

Note: A harder problem than Mountain Car due to underactuation and nonlinear dynamics.

Modifications to Classic CartPole for Swing-Up Task

- **Initial Condition:**

- Pole starts hanging downward — requires active swing-up to reach upright.

- **Reward Function Modified:**

- Reward is based on `upright = cos()`, giving:
 - Maximum reward when pole is upright ($\theta = 0$)
 - Zero reward when hanging down ($\theta = \pi$)

- **No Early Termination:**

- Classic CartPole ends episode when angle deviates too much.
- The wrapper ensures episode ends only when cart moves beyond threshold.

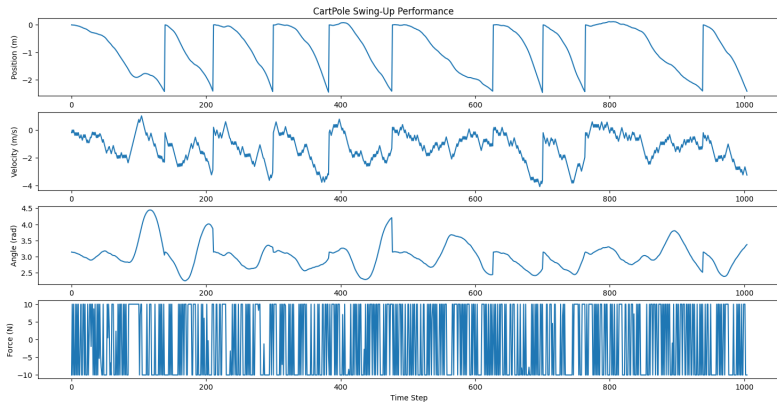
- **Continuous Actions and Force Scaling:**

- Supports continuous control input.
- Action made continuous clipped to range: $[-10, 10]$

Adapting the Algorithm for Swing-Up Pendulum

- **Modified Neural Network Architecture:**
 - Input dimension increased to 4: $(x, \dot{x}, \theta, \dot{\theta})$
 - Network expanded to hidden layer with **256 neurons**
- **Increased Exploration Steps:**
 - More episodes (100) run with random actions for better coverage of state space
- **Extended iLQR Planning Horizon:**
 - Longer horizon used to account for the swing-up and balance phases
- **More iLQR Optimization Iterations:**
 - Increased inner loop iterations for better trajectory refinement

Results: ILQR + NN base implementation



Very short runs, sub optimal result

Core Technical Improvements

1. Experience Replay Buffer

- Stores 100K transitions
- Batched training (64 samples)
- Breaks temporal correlation

2. PyTorch Dynamics Model

- Automatic differentiation
- GPU acceleration support
- Precise Jacobians via autograd

3. Residual Network

- Output of NN is difference from the previous state [Delta x].
- Helps converge faster and give more smoother dynamics.

4. Enhanced Architecture

- 256 hidden units
- ReLU activation
- ADAM optimizer
- MSE loss

5. Training Protocol

- 1 Collect episode data
- 2 Train model (50 epochs)
- 3 Optimize with iLQR
- 4 Repeat

6. Adaptive Line Search


$$u_t^{\text{new}} = u_t + \alpha k_t + K_t(x_t - x_t^{\text{ref}})$$

- Guarantees cost reduction
- Backtracking with $\alpha = 0.5 \rightarrow 0.7$


7. Regularized Cost Function

$$J = \underbrace{10\theta^2}_{\text{Upright}} + \underbrace{0.1x^2}_{\text{Position}} + \underbrace{0.1v^2}_{\dot{x}} + \underbrace{0.1x^2}_{\dot{\theta}}$$



Comparison of Results from Different Runs



Run 1



Run 2



Conclusion

- Successfully implemented **iLQR** to solve classic control tasks in the Gymnasium framework.
- Demonstrated effective use of a **neural network-based dynamics model** to handle unknown system dynamics.
- Applied the method to both **Mountain Car** and the more challenging **Swing-Up CartPole** environment.
- Incorporated key adaptations: deeper neural networks, extended horizons, and enhanced iLQR optimization to improve performance on complex tasks.

Key Takeaway

iLQR combined with learned dynamics can serve as a powerful model-based control strategy even in challenging nonlinear settings.

- Learn unknown parameters like mass and gravity explicitly.
- Extend to more complex environments.
- Improve sample efficiency of the training loop.

Thank You!

Questions and Discussion Welcome.